




Command: `aprun -n 64 ./wave_cray`  
Resources: 2 nodes (32 physical, 64 logical cores per node)  
Memory: 63 GiB per node  
Tasks: 64 processes  
Machine: `beskow-login1.pdc.kth.se`  
Start time: Thu Sep 8 14:57:38 2016  
Total time: 31 seconds  
Full path: `/cfs/klemming/scratch/c/cira/students/allinea-reports-6.1`



## Summary: wave\_cray is **Compute-bound** in this configuration

**Compute** 77.4%

Time spent running application code. High values are usually good. This is **high**; check the CPU performance section for advice.

**MPI** 22.6%

Time spent in MPI calls. High values are usually bad. This is **low**; this code may benefit from a higher process count.

**I/O** 0.0%

Time spent in filesystem I/O. High values are usually bad. This is **negligible**; there's no need to investigate I/O performance.

This application run was **Compute-bound**. A breakdown of this time and advice for investigating further is in the **CPU** section below. As little time is spent in **MPI** calls, this code may also benefit from running at larger scales.

### CPU

A breakdown of the 77.4% CPU time:

Scalar numeric ops 5.3%  
Vector numeric ops 14.7%  
Memory accesses 80.0%

The per-core performance is **memory-bound**. Use a profiler to identify time-consuming loops and check their cache performance.

### MPI

A breakdown of the 22.6% MPI time:

Time in collective calls 2.8%  
Time in point-to-point calls 97.2%  
Effective process collective rate 1.13 MB/s  
Effective process point-to-point rate 2.64 MB/s

Most of the time is spent in **point-to-point calls** with a **very low** transfer rate. This suggests load imbalance is causing synchronization overhead; use an MPI profiler to investigate.

### I/O

A breakdown of the 0.0% I/O time:

Time in reads 0.0%  
Time in writes 0.0%  
Effective process read rate 0.00 bytes/s  
Effective process write rate 0.00 bytes/s

No time is spent in **I/O** operations. There's nothing to optimize here!

### Threads

A breakdown of how multiple threads were used:

Computation 0.0%  
Synchronization 0.0%  
Physical core utilization 99.3%  
System load 100.0%

No measurable time is spent in multithreaded code.

### Memory

Per-process memory usage may also affect scaling:

Mean process memory usage 97.2 MiB  
Peak process memory usage 108 MiB  
Peak node memory usage 11.0%

The **peak node memory usage** is very low. Running with fewer MPI processes and more data on each process may be more efficient.